

# Semantic Image Inpainting with Deep Generative Models

Raymond A. Yeh\*, Chen Chen\*, Teck Yian Lim,  
Alexander G. Schwing, Mark Hasegawa-Johnson, Minh N. Do  
University of Illinois at Urbana-Champaign

{yeh17, cchen156, tlim11, aschwing, jhasegaw, minhdo}@illinois.edu

*CVPR 2017*

Slides compiled by Lars Gjesteby  
July 5, 2018

# Motivation

- Estimate missing image data
- Overcome limitations of existing methods that use information only from a single image

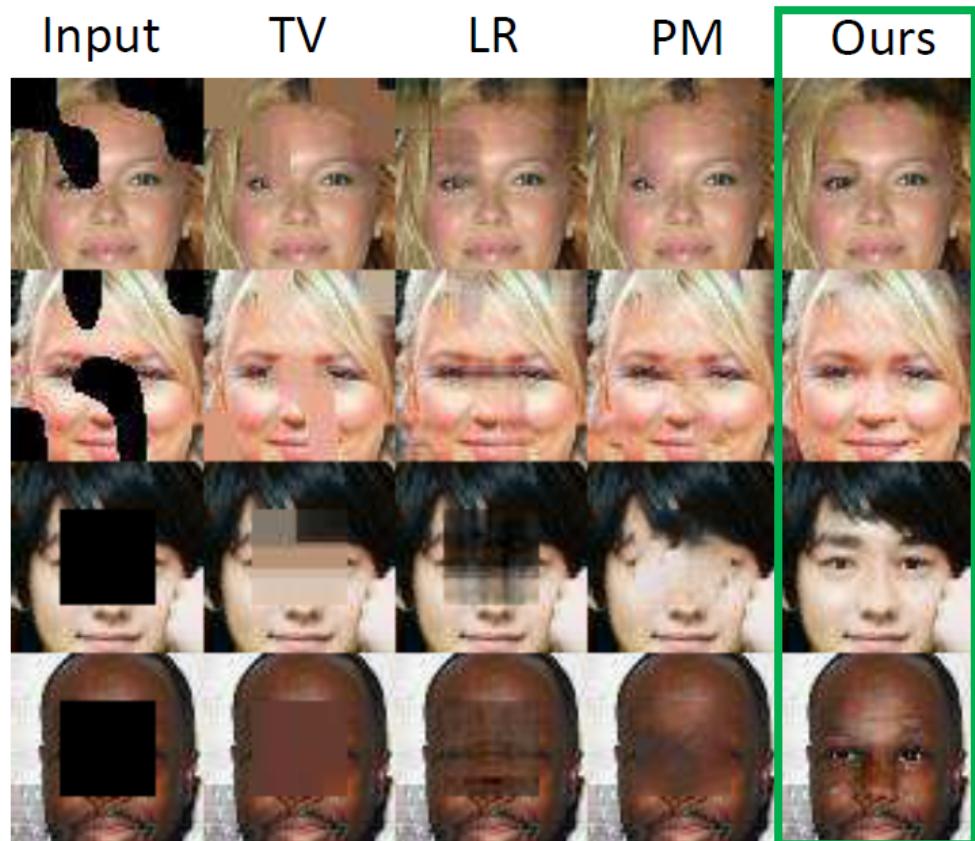
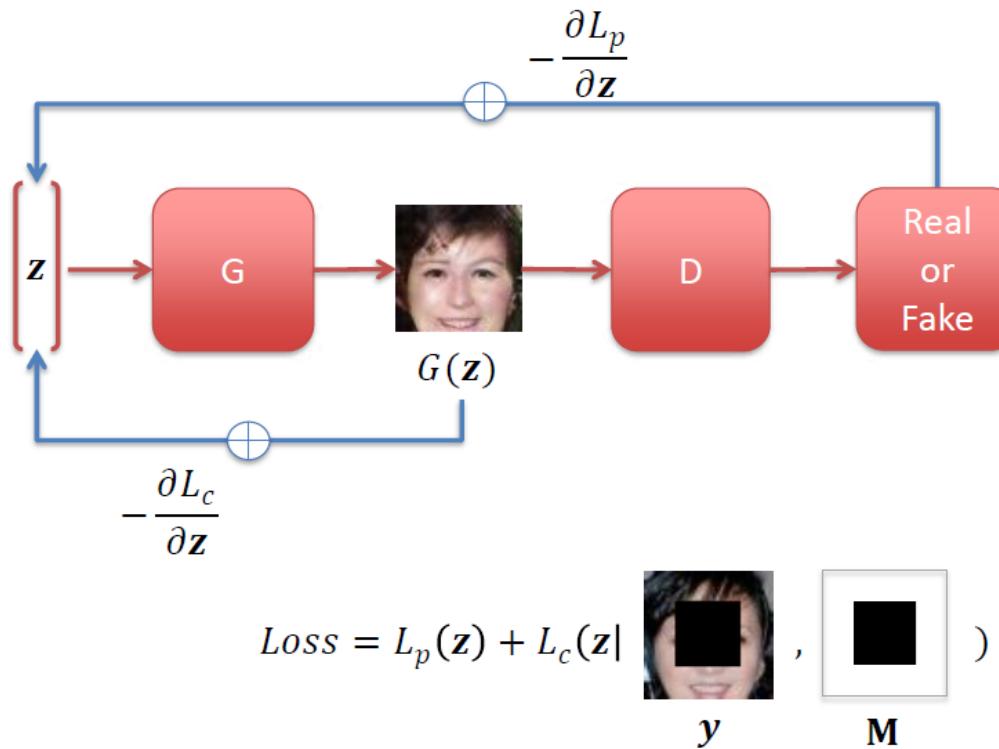


Figure 1. Semantic inpainting results by TV, LR, PM and our method. Holes are marked by black color.

TV – Total variation  
LR – Low rank  
PM – PatchMatch

# Approach

- Semantic inpainting with a generative adversarial network model (DCGAN from Radford *et al.*)
- Train generator on uncorrupted data



$p_{\mathbf{z}}$  – corrupt data distribution  
 $p_{\text{data}}$  – true data distribution  
 $\mathbf{z}$  – learned encoding manifold  
 $\mathbf{y}$  – corrupted image  
 $\mathbf{M}$  – binary mask

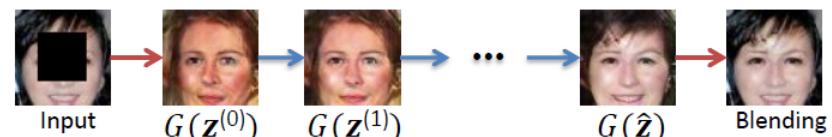
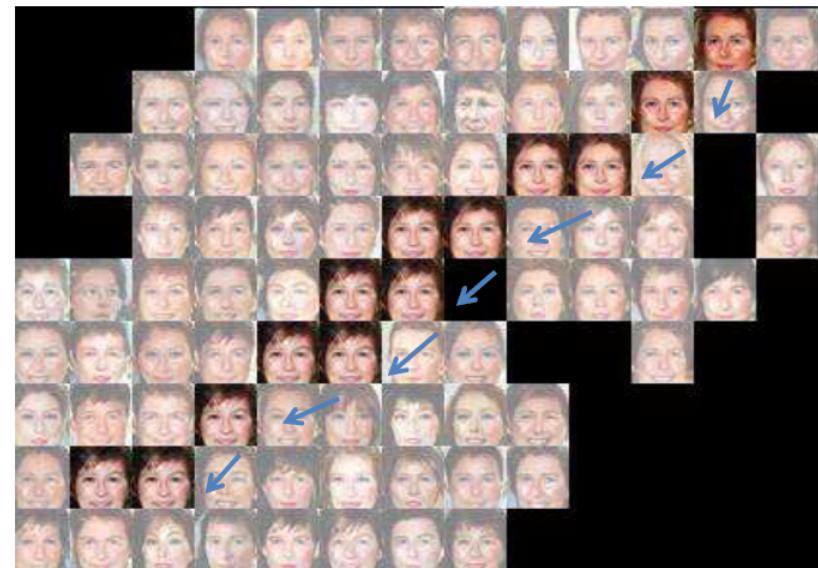
# Approach

- Training  $G$  to be efficient – an image not from  $p_{data}$  should not lie on  $\mathbf{z}$
- Iteratively update  $\mathbf{z}$  to find closest mapping to corrupt image while constrained to manifold ( $\hat{\mathbf{z}}$ )

$$\hat{\mathbf{z}} = \arg \min_{\mathbf{z}} \{ \mathcal{L}_c(\mathbf{z} | \mathbf{y}, \mathbf{M}) + \mathcal{L}_p(\mathbf{z}) \}$$

$\mathcal{L}_c$  – context loss: constrains generated image to input  $\mathbf{y}$  and mask  $\mathbf{M}$

$\mathcal{L}_p$  – prior loss: penalizes unrealistic images



# Loss Functions

- Context loss: a weighted  $\ell_1$ -norm difference between the recovered image and the uncorrupted portion

$$\mathcal{L}_c(\mathbf{z}|\mathbf{y}, \mathbf{M}) = \|\mathbf{W} \odot (G(\mathbf{z}) - \mathbf{y})\|_1$$

$$\mathbf{W}_i = \begin{cases} \sum_{j \in N(i)} \frac{(1-\mathbf{M}_j)}{|N(i)|} & \text{if } \mathbf{M}_i \neq 0 \\ 0 & \text{if } \mathbf{M}_i = 0 \end{cases}$$

$N(i)$  – pixel neighbors in a local window

- Prior loss: encourages recovered image to be similar to the samples from training set

$$\mathcal{L}_p(\mathbf{z}) = \lambda \log(1 - D(G(\mathbf{z})))$$

$\lambda$  – balancing parameter

# Effect of Prior Loss

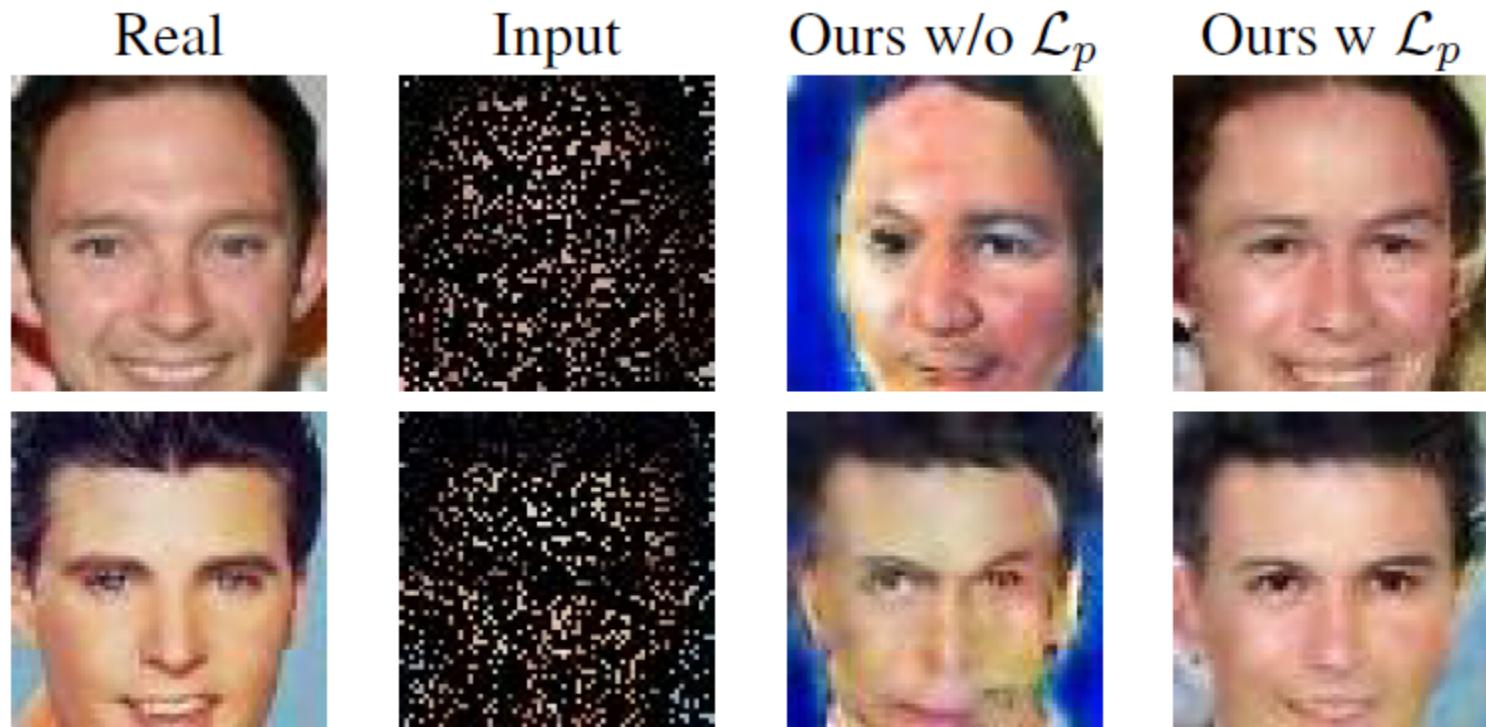


Figure 4. Inpainting with and without the prior loss.

# Poisson Blending

- Predicted pixels may not exactly preserve the same intensities of the surrounding pixels, although the content is correct and well aligned
- Keep gradients of  $G(\hat{\mathbf{z}})$  to preserve image details while shifting the color to match the color in  $\mathbf{y}$

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \|\nabla \mathbf{x} - \nabla G(\hat{\mathbf{z}})\|_2^2,$$

$$\text{s.t. } \mathbf{x}_i = \mathbf{y}_i \text{ for } \mathbf{M}_i = 1,$$

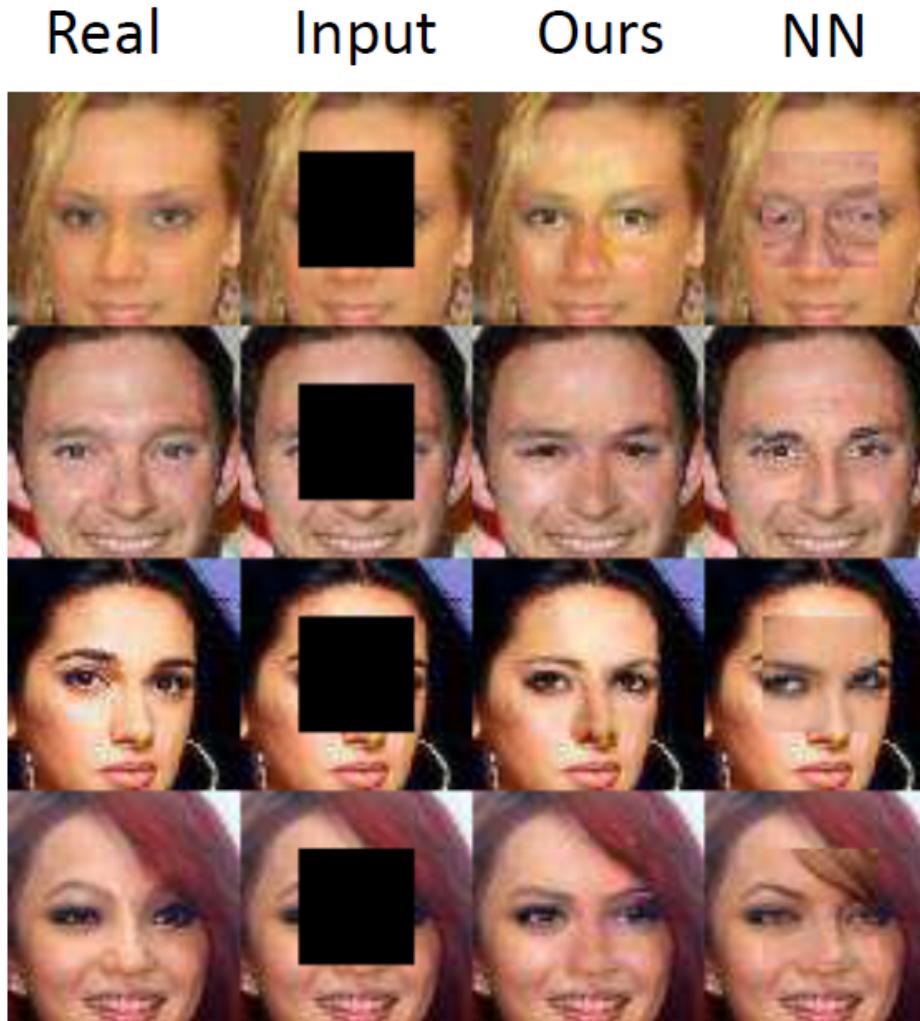


Figure 5. Inpainting with and without blending.

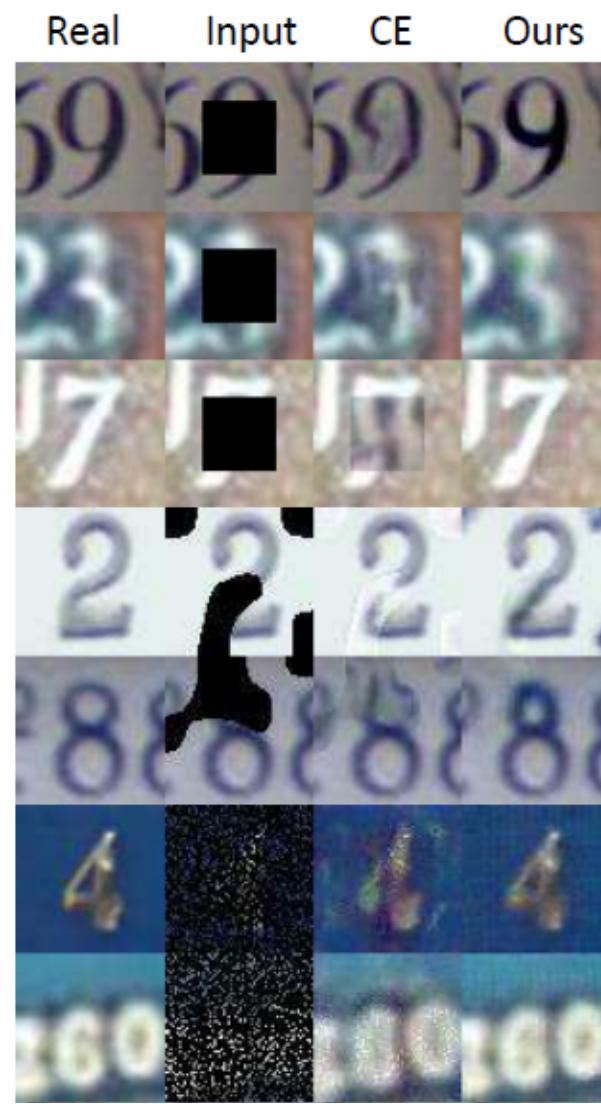
# Validation

- Datasets:
  - CelebFaces Attributes (CelebA)
  - Street View House Numbers (SVHN)
  - Stanford Cars
- Four masks:
  1. Central block
  2. Random pattern masks, with approximately 25% missing
  3. 80% missing complete random masks
  4. Half missing masks (randomly horizontal or vertical)

# Compare with Nearest Neighbor



# Compare with Context Encoder



# Quantitative Results

Table 1. The PSNR values (dB) on the test sets. Left/right results are by CE[30]/ours.

| Masks/Dataset | CelebA            | SVHN              | Cars              |
|---------------|-------------------|-------------------|-------------------|
| Center        | <b>21.3</b> /19.4 | <b>22.3</b> /19.0 | <b>14.1</b> /13.5 |
| pattern       | <b>19.2</b> /17.4 | <b>22.3</b> /19.8 | 14.0/ <b>14.1</b> |
| random        | 20.6/ <b>22.8</b> | 24.1/ <b>33.0</b> | 16.1/ <b>18.9</b> |
| half          | <b>15.5</b> /13.7 | <b>19.1</b> /14.6 | <b>12.6</b> /11.1 |

# Paper Conclusions

- Proposed model provides better visual image quality over state-of-the-art methods
- Works well for relatively simple structures like faces, but is too small to represent complex scenes in the world
- Prediction performance strongly relies on the generative model and the training procedure



Figure 13. Some failure examples.